

# ParallelPrefix

March 19, 2019

## 1 Parallel Prefix

Let's try to write first the sequential version of the parallel prefix. Initially we will just consider a short vector, than we'll assemble the actual main code with larger vectors. First we define the standard pre-amble

```
In [1]: #include <iostream>
        #include <vector>

        using namespace std;
```

We define the sample vector to compute

```
In [2]: vector<int> y = {1,2,3,4,5,6,7,8};
```

and the sequential code, needed to take into account the “quality” of what we are going to parallelize ...

```
In [3]: int oplus(int x, int y) {
        auto temp = x+y;
        // possibly we will fill here some delay/longer computation
        return temp;
    }
```

```
In [4]: void prefix(vector<int> &x, function<int(int,int)> oplus) {
        for(int i=1; i<x.size(); i++)
            x[i]= oplus(x[i],x[i-1]);
        return;
    }
```

```
In [6]: void printv(vector<int> x){
        for(int i=0; i<x.size(); i++)
            cout << x[i] << " ";
        cout << endl;
        return;
    }
```

```

input_line_13:1:6: error: redefinition of 'printf'
void printf(vector<int> x){
    ^
input_line_12:1:6: note: previous definition is here
void printf(vector<int> x){
    ^

```

```
In [7]: printf(y);
```

```
1 2 3 4 5 6 7 8
```

Now we can compute the prefix and see if the vector has changed ...

```
In [8]: prefix(y,oplus);
```

```
In [9]: printf(y);
```

```
1 3 6 10 15 21 28 36
```

which is ok. In order to have an idea of what's going on here, let's try to measure how much time we spend computing the prefix. We first define the usual timer:

```
In [10]: #include <chrono>
         #include <thread>
```

```

class utimer {
    std::chrono::system_clock::time_point start;
    std::chrono::system_clock::time_point stop;
    std::string message;
    using usecs = std::chrono::microseconds;
    using msec = std::chrono::milliseconds;

public:

    utimer(const std::string m) : message(m) {
        start = std::chrono::system_clock::now();
    }

    ~utimer() {
        stop =
            std::chrono::system_clock::now();
        std::chrono::duration<double> elapsed =
            stop - start;
        auto msec =
            std::chrono::duration_cast<std::chrono::microseconds>(elapsed).count();
    }

```

```
std::cout << message << " computed in " << msec << " usec "
<< std::endl;

}
};

In [11]: vector<int> x = {1,2,3,4,5,6,7,8};

{
    utimer t("prefix8");
    prefix(x, oplus);
}

prefix8 computed in 8 usec
```

```
In [12]: vector<int> longx(1024);
         for(auto &i : longx) i = 1;
```

[illegible]

longx computed in 1537 usec

1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35

```
In [16]: void active_udelay(int usecs) {
    // read current time
    auto start = std::chrono::high_resolution_clock::now();
    auto end   = false;
    while(!end) {
        auto elapsed = std::chrono::high_resolution_clock::now() - start;
```

```

        auto usec = std::chrono::duration_cast<std::chrono::microseconds>(elapsed).count();
        if(usec>usecs)
            end = true;
    }
    return;
}

```

```

In [17]: int oplus2(int x, int y) {
        auto temp = x + y;
        active_udelay(1000);
        return temp;
    }

```

```

In [ ]: cout << oplus(1,1) << endl;

```

```

In [18]: vector<int> w = {1,1,1,1,1,1,1,1};

```

```

In [19]: prefix(w,oplus2);

```

```

In [20]: printv(w);

```

```

1 2 3 4 5 6 7 8

```

```

In [21]: {
        utimer t3("now");
        prefix(w,oplus2);
    }

```

```

now computed in 7044 usec

```

```

In [ ]: const int init = 1;
        vector<int> d(1024);

```

```

In [ ]: for(auto &i: d) i=2;

```

```

In [ ]: printv(d);

```

```

In [ ]: {
        utimer t("longy");
        prefix(d, oplus2);
    }

```

```

In [ ]: // printv(d);

```

## 2 Parallel version

We now define the parallel version. We assume the length of the vector is a multiple of the parallelism degree. We compute the parallel prefix by: 1) computing the local prefix on each partition 2) computing the prefix of the vector made of the last item computed in each partition 3) adding in partition  $i$  (excluding the first one) the sum of current item with the  $i-1$  position of the prefix vector just computed

```
In [ ]: printv(w);
```

In the first phase, we have to compute the local prefixes:

```
In [22]: void phase1(vector<int> &x,          // the input/output vector
                    vector<int> &p, // the prefix vector
                    int i,          // the partition index
                    function<int(int,int)> oplus, // the combiner function
                    int nw         // the parallelism degree
                    ) {

    auto m = x.size();
    auto delta = m / nw;
    // phase one, compute prefix of the assigned partition
    for(int j=i*delta+1; j<=(i+1)*delta; j++)
        x[j] = oplus(x[j-1], x[j]);
    // and assign final item to the prefix vector
    p[i] = x[(i+1)*delta-1];
    // cout << "OK";
    return;
}
```

Assuming we have 2 workers, we compute the vector  $p$  of prefixes:

```
In [23]: printv(w);
```

```
1 3 6 10 15 21 28 36
```

```
In [24]: vector<int> p(2);
```

```
In [25]: vector<int> q = {1,2,3,4,5,6,7,8};
```

```
In [27]: phase1(q,p,0,oplus,2);
```

```
In [28]: printv(q);
```

```
1 3 6 10 5 6 7 8
```

```
In [29]: printv(p);
```

10 0

```
In [30]: phase1(q,p,1,oplus,2);
```

```
In [31]: printv(p);
```

10 26

```
In [32]: printv(q);
```

1 3 6 10 5 11 18 26

Then we have to compute the prefix of the prefixes

```
In [33]: prefix(p,oplus);
```

```
In [34]: printv(p);
```

10 36

```
In [35]: printv(q);
```

1 3 6 10 5 11 18 26

and eventually we have to use the vector to update the partitions

```
In [36]: void phase2(vector<int> &x, vector<int> &p, int i, function<int(int,int)> opus, int nw)
        {
            auto m = x.size();
            auto delta = m /nw;
            for(int j=i*delta; j<(i+1)*delta; j++)
                x[j] = opus(x[j], p[i-1]);
            return;
        }
```

the first partition does not need to be updated ...

```
In [ ]: // phase2(w,p,0,oplus,2);
```

the rest of the partitions need to be updated

```
In [38]: phase2(q,p,1,oplus,2);
```

```
In [39]: printv(q);
```

1 3 6 10 15 21 28 36

Now we can arrange things in parallel. The naive solution creates and join threads for both phases:

```
In [ ]: auto f = [] (int x) { cout << x << endl; return; };
```

```
In [ ]: void foo()
{
    // simulate expensive operation
    std::this_thread::sleep_for(std::chrono::seconds(1));
    cout << "ciao" << endl;
}
```

```
In [ ]: #include <thread>
        std::thread helper2(foo);
```

```
In [ ]: helper2.join();
```

```
In [40]: void compute_prefix(vector<int> &x, function<int(int,int)> oplus, int nw) {

    vector<thread*> tid(nw);
    vector<int> p(nw);

    // first create threads for phase 1
    for(int i=0; i<nw; i++)
        tid[i] = new thread(phase1, ref(x), ref(p), i, oplus, nw);
    for(int i=0; i<nw; i++)
        tid[i]->join();

    // then compute prefix of prefixes
    // prefix(p,ref(oplus));
    for(int i=1; i<nw; i++)
        p[i] = oplus(p[i-1],p[i]);

    // then threads for phase 2
    for(int i=0; i<nw; i++)
        tid[i] = new thread(phase2, ref(x), ref(p), i, oplus, nw);
    for(int i=0; i<nw; i++)
        tid[i]->join();
    return;
}
```

```
In [41]: vector<int> a(8,1);
```

```
In [42]: printv(a);
```

```
1 1 1 1 1 1 1 1
```

```
In [43]: compute_prefix(a,oplus,2);
```

```
In [44]: printv(a);
```

1 2 3 4 5 6 7 8

this works, let's add some timing ...

```
In [45]: void compute_prefix_timed(vector<int> &x, function<int(int,int)> oplus, int nw) {
```

```

vector<thread*> tid(nw);
vector<int> p(nw);

{
    utimer t0("all");

    // first create threads for phase 1
    {
        utimer t1("phase1");
        for(int i=0; i<nw; i++)
            tid[i] = new thread(phase1, ref(x), ref(p), i, oplus, nw);
        for(int i=0; i<nw; i++)
            tid[i]->join();
    }

    // then compute prefix of prefixes
    // prefix(p,ref(oplus));
    {
        utimer t1("phase2");
        for(int i=1; i<nw; i++)
            p[i] = oplus(p[i-1],p[i]);
    }

    // then threads for phase 2
    {
        utimer t1("Phase3");
        for(int i=0; i<nw; i++)
            tid[i] = new thread(phase2, ref(x), ref(p), i, oplus, nw);
        for(int i=0; i<nw; i++)
            tid[i]->join();
    }
}

return;
}

```

```
In [46]: vector<int> b(1024,1);
```

```
In [47]: printv(b);
```

[illegible]



```

In [48]: compute_prefix_timed(b,oplus,4);

phase1 computed in 718 usec
phase2 computed in 5 usec
Phase3 computed in 1341 usec
all computed in 2479 usec


In [ ]: vector<int> c(1024,1);

In [ ]: compute_prefix_timed(b,oplus,1);

In [ ]: vector<int> a1(1024,1);

In [ ]: vector<int> a2(1024,1);

In [ ]: compute_prefix_timed(a1,oplus2,4);

In [ ]: compute_prefix_timed(a2,oplus2,1);

In [ ]: {
    utimer("seq");

    vector<int> a1(1024,1);
    for(int i=1; i<1024; i++)
        a1[i] = oplus2(x[i-1],x[i]);

}

In [ ]: cout <<
        oplus2(1,2);

```